

Deep Learning Project 1

MLPs and CNN for Image Classification

Mavroforos Nikolaos - p3352014

Ntenezos Panagiotis - p3352025

March 9, 2022

The code can be found here: <https://github.com/d-tchmnt/Deep-Learning>

1 Introduction

Given an image of a fashion item, build a deep learning model that recognizes the fashion item. You must use at least 2 different architectures, one with MLPs and one with CNNs. Use the Fashion-MNIST dataset to train and evaluate your models. More information about the task and the dataset can be found at <https://github.com/zalandoresearch/fashion-mnist>. The dataset is also available from Tensorflow and Keras.

2 Data Manipulation

2.1 Dataset

The dataset, after retrieving it from Keras, contains 60,000 28×28 grayscale images of images associated with a label from 10 fashion categories, along with a test set of 10,000 images. Each Label corresponds to the integer of the class to which the item belongs.

2.2 Data split

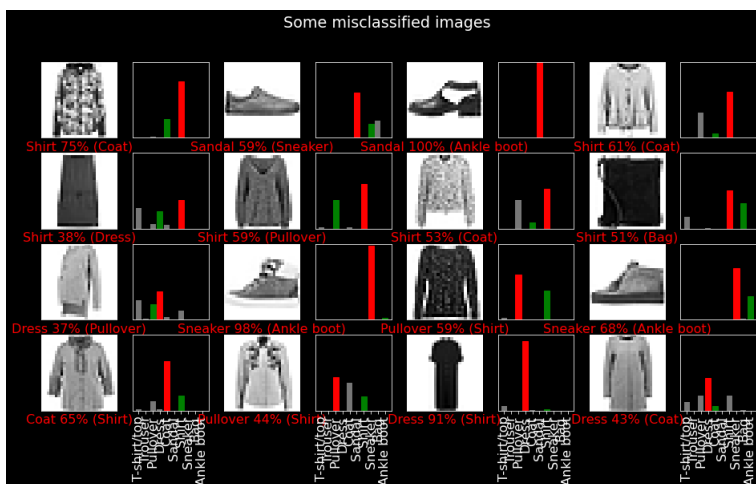
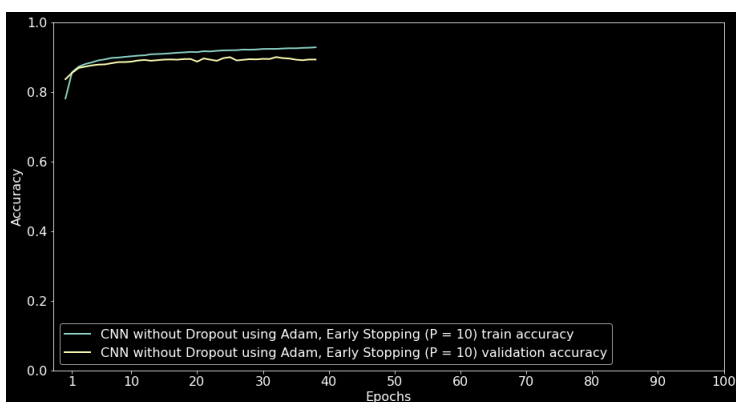
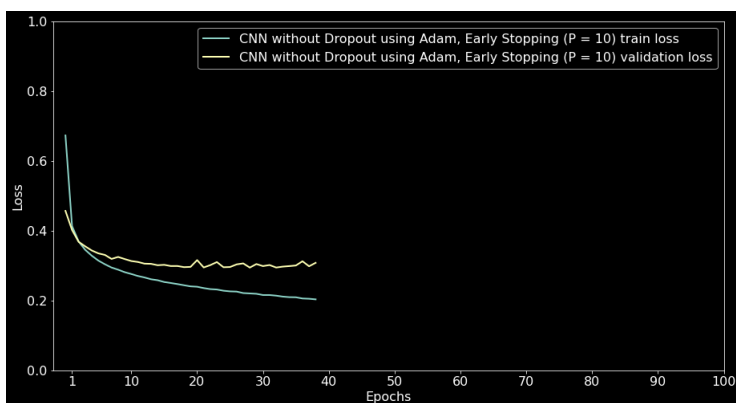
We use the `train_test_split` the dataset into train, and dev train while retrieving the test dataset directly from the keras embedding dataset. We used the `stratified` option, while also making sure `shuffle` is enabled. This way, we shuffle the data before we split them (useful in case data are pre-arranged or sorted) while making sure that the generated subsets of data contain classes with the distribution.

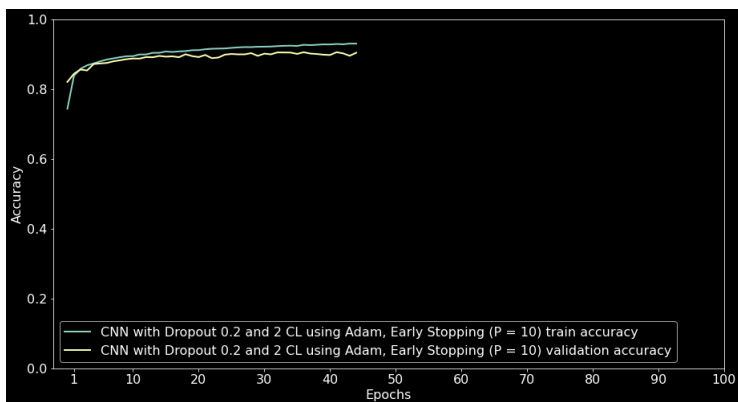
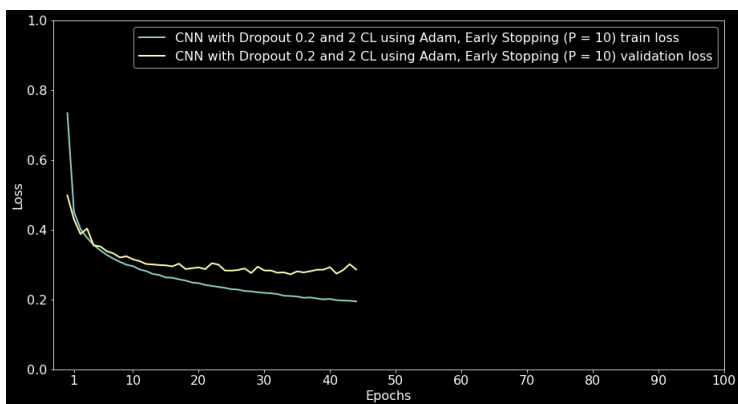
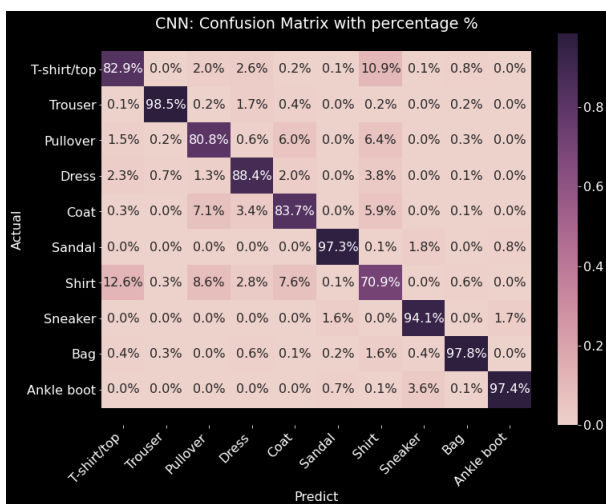
Then we normalize the input pixel values, by dividing each input by 255 (the number of pixels per image) so that every image has embedded range from 0 to 1. This is useful for making faster calculations since Neural networks process inputs using small weight values, and inputs with large integer values can disrupt or slow down the learning process. As such it is good practice to normalize the pixel values so that each pixel value has a value between 0 and 1.

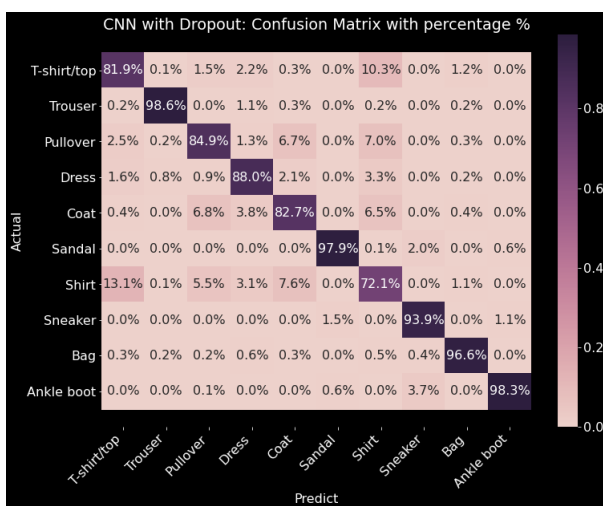
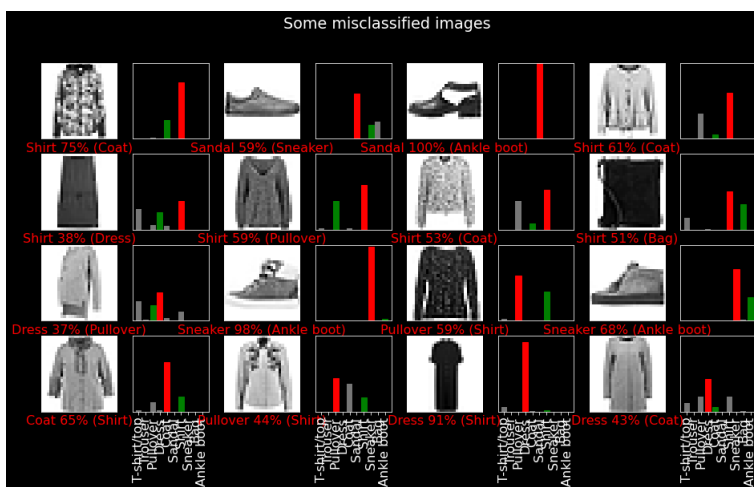
2.3 Data Exploratory Analysis

3 MLP Classifier

4 CNN Classifier







Train Loss : 0.24073

Validation Loss: 0.24195

Test Loss : 0.24059

Train Accuracy : 0.91205

Validation Accuracy: 0.90745

Test Accuracy : 0.91250

